

1. How many bytes would the following array declaration allocate on a 64-bit machine?

```
char *arr[10][6];
```

480 bytes. Each char pointer is 8 bytes. There are 10 rows by 6 columns in the 2D array. $8 \times 10 \times 6 = 480$.

2. What will the following print out?

```
typedef struct {
    char shookie;
    int tata;
    char cookie;
    double chimmy;
} bt;

void main(int argc, char** argv){
    bt band[7];
    printf( "%d\n", (int)sizeof(band));
}
```

$(1 + (3) + 4 + 1 + (7) + 8) * 7 = 168$

Due to alignment, we need to add the numbers in parentheses

3. What is the best* ordering of the following variables if you want to have a struct that uses all of them? Assume a 64-bit architecture with 4-byte ints.

** the ordering that will result in the optimal usage of space.*

```
char tully;
long stark;
float* lannister;
double targaryen;
int greyjoy;
float arryn;
```

```
// want to order from largest size to smallest, as structs are
// x-aligned, where x is the size of the largest data type in
// the struct
```

```
struct Westeros{
    float* lannister;    // ALL pointers are 8 bytes
    double targaryen;   // doubles are 8 bytes
    long stark;         // longs are 8 bytes
    float arryn;        // floats are 4 bytes
    int greyjoy;        // ints are 4 bytes
    char tully;         // chars are 1 byte
};
```

4. Consider the following disassembled function:

```
000000000040102b <phase_2>:
  40102b:  55                push    %rbp
  40102c:  53                push    %rbx
  40102d:  48 83 ec 28       sub     $0x28,%rsp
  401031:  48 89 e6          mov     %rsp,%rsi
  401034:  e8 e3 03 00 00    callq   40141c
<read_six_numbers>
  401039:  83 3c 24 01       cmpl    $0x1, (%rsp)
  ...
```

Right after the callq instruction has been executed, what address will be at the top of the stack?

401039.

- When executing a call instruction, you push the return address onto the stack
 - The instruction pointer (%rip) points to the next instruction to execute
 - In this case, 401039
- When you reach the ret instruction in read_six_numbers, you will pop this address off the stack so control will return to the next instruction in phase_2.

5. Consider the following C code:

```
typedef struct {
    char first;
    int second;
    short third;
    int* fourth;
} stuff;

stuff array[5];

int func0(int index, int pos, long dist) {
    char* ptr = (char*) &(array[index].first);
    ptr += pos;
    *ptr = index + dist;

    return *ptr;
}

int func1() {
    int x = func0(1, 4, 12);
    return x;
}
```

Clearly some code is missing - your job is to fill in the blanks! Note that the size of the blanks is not significant. The two functions will be compiled using the following assembly code:

```
0000000000400492 <func0>:
 400492: 8d 04 17                lea    (%rdi,%rdx,1),%eax
 400495: 48 63 ff                movslq %edi,%rdi
 400498: 48 63 f6                movslq %esi,%rsi
 40049b: 48 8d 14 7f            lea    (%rdi,%rdi,2),%rdx
 40049f: 88 84 d6 60 10 60 00    mov    %al,0x601060(%rsi,%rdx,8)
 4004a6: 0f be c0                movsbl %al,%eax
 4004a9: c3                      retq

00000000004004aa <func1>:
 4004aa: c6 05 cb 0b 20 00 0d    movb    $0xd,0x200bcb(%rip)
                                # 60107c <array+0x1c>
 4004b1: b8 0d 00 00 00          mov     $0xd,%eax
```

4004b6: c3

retq

The answer can be derived by tackling func0 first, then func1
func0

- From instruction 400492, we can see that the return value is set to %rdi + %rdx, where %rdi is index and %rdx is dist
 - %rdi is set to the first parameter, %rsi to the second parameter, %rdx to the third
 - %eax is unchanged, until instruction 4004a6 with %al
 - This makes sense, since we're returning the value from dereferencing a pointer to a char, aka a single byte (%al is a single byte)
 - Thus we know ***ptr = index + dist**
- From instruction 40049b:
 - %rdx is set to 3 * %rdi
 - %rdx is thus 3 * index
- From instruction 40049f:
 - 0x601060 is presumably the start of the array
 - This is confirmed in instruction 4004aa, where 60107c is shown to be <array+0x1c>
 - The destination of instruction 40049f is thus:
 - (Start of the array) +
8 * (3 * %rdi) +
pos
 - = (start of array) + (24 * index) + pos
 - Each object of type stuff is 24 bytes (alignment)
 - ptr from func0 is thus pointing to **array[index].first**
 - The "+ pos" comes from the second line of func0

func1

- **(note) there is no call to func0, as this code was produced from gcc -O**
 - Optimization has not been covered yet, but the gist of the problem should be understandable
- From Week3 Lecture slides "data_examples.pdf", students should understand that 0x200bcb(%rip) from instruction 4004aa is location <array + 0x1c>
 - 0x1c = 28
 - Since each object of type stuff is 24 bytes, we know the second parameter (pos) was called with value 4
 - array[1].first would be at byte 24
 - ptr += 4 would bring us to 28
 - Thus we know **pos = 28 - 24 = 4**

- 0xd = 13
 - Thus we know that the **third parameter (dist) was called with value 12**